

Sep 07, 22 23:42	main.py	Page 1/4
<pre>#!/bin/env python3.8 """ Simon Yoon ECE470 Deep Learning Professor Curro Took inspiration from Professor Curro's example submission Husam Almanakly and Michael Bentivegna helped debug; There were issues in methods that I used both tf and np simultaneously in. Changing it so that it was one or the either locally fixed it --believed it to be a datatype handling issue i.e tf.zeros vs np.zeros, tf.exp vs np.exp """ import os import logging import matplotlib import matplotlib.pyplot as plt import numpy as np import tensorflow as tf from absl import app from absl import flags from tqdm import trange from dataclasses import dataclass, field, InitVar script_path = os.path.dirname(os.path.realpath(__file__)) @dataclass class LinearModel: weights: np.ndarray bias: float m: np.ndarray # mu s: np.ndarray # sigma @dataclass class Data: model: LinearModel rng: InitVar[np.random.Generator] num_features: int num_samples: int sigma: float x: np.ndarray = field(init=False) y: np.ndarray = field(init=False) def __post_init__(self, rng): self.index = np.arange(self.num_samples) self.x = rng.uniform(0.1, 0.9, size=(self.num_samples, 1)) # self.num_features -> 1 clean_y = np.sin(2 * np.pi * self.x) # + epsilon e = rng.normal(loc=0, scale=0.2, size=(self.num_samples, 1)) # e_i is drawn from N(0, sigma_noise) something close to 0.1 self.y = clean_y + e # y_i = sin(2*pi*x_i)+e</pre>		

Sep 07, 22 23:42	main.py	Page 2/4
<pre> def get_batch(self, rng, batch_size): """ Select random subset of examples for training batch """ choices = rng.choice(self.index, size=batch_size) return self.x[choices], self.y[choices].flatten() """ def compare_linear_models(a: LinearModel, b: LinearModel): for w_a, w_b in zip(a.weights, b.weights): print(f"{w_a:0.2f}, {w_b:0.2f}") print(f"{a.bias:0.2f}, {b.bias:0.2f}") """ font = { # "family": "Adobe Caslon Pro", "size": 10, } matplotlib.style.use("classic") matplotlib.rc("font", **font) FLAGS = flags.FLAGS flags.DEFINE_integer("num_features", 6, "Number of bases in record") # M = 6 flags.DEFINE_integer("num_samples", 50, "Number of samples in dataset") # N = 50 flags.DEFINE_integer("batch_size", 16, "Number of samples in batch") flags.DEFINE_integer("num_iters", 3000, "Number of SGD iterations") flags.DEFINE_float("learning_rate", 0.05, "Learning rate / step size for SGD") # 0.1 -> 0.05 for oscillation on smaller magnitude iters flags.DEFINE_integer("random_seed", 31415, "Random seed") flags.DEFINE_float("sigma_noise", 0.1, "Standard deviation of noise random variable") # 0.5 -> 0.1 flags.DEFINE_bool("debug", False, "Set logging level to debug") class Model(tf.Module): def __init__(self, rng, num_features): """ A plain linear regression model with a bias term altered to be a regression model to estimate sinusoid with weights (w), bias (b), mu (m), sigma (s) """ self.num_features = num_features # trainable params self.w = tf.Variable(rng.normal(shape=[self.num_features]), name="weights") self.b = tf.Variable(tf.zeros(shape=[1, 1]), name="bias") self.m = tf.Variable(rng.normal(shape=[self.num_features]), name="mean") # mu self.s = tf.Variable(rng.normal(shape=[self.num_features]), name="sigma") # stdev def __call__(self, x): """ Functional form with expanded phi </pre>		

Sep 07, 22 23:42

main.py

Page 3/4

```

Gaussian basis functions
"""
    gau = tf.zeros(shape=(x.shape[0], 1))
    """
for i in range(self.num_features):
    gau += tf.math.exp(-1 * (((x - self.m[i]) ** 2) / (self.s[i] ** 2)))
gau += self.b
"""
    gau = self.w * tf.math.exp(-1 * (((x - self.m) ** 2) / (self.s**2)))
    return tf.squeeze(tf.reduce_sum(gau, 1) + self.b)  # incl mu, sig est

@property
def model(self):
    return LinearModel(
        self.m.numpy().reshape([self.num_features]),
        self.s.numpy().reshape([self.num_features]),
        self.w.numpy().reshape([self.num_features]),
        self.b.numpy().squeeze(),
    )

def main(a):
    logging.basicConfig()

    if FLAGS.debug:
        logging.getLogger().setLevel(logging.DEBUG)

    # Safe np and tf PRNG
    seed_sequence = np.random.SeedSequence(FLAGS.random_seed)
    np_seed, tf_seed = seed_sequence.spawn(2)
    np_rng = np.random.default_rng(np_seed)
    tf_rng = tf.random.Generator.from_seed(tf_seed.entropy)

    data_generating_model = LinearModel(
        weights=np_rng.integers(low=0, high=5, size=(FLAGS.num_features)),
        bias=2,
        m=np_rng.integers(low=0, high=1, size=(FLAGS.num_features)),
        s=np_rng.integers(low=0, high=1, size=(FLAGS.num_features)),
    )
    logging.debug(data_generating_model)

    data = Data(
        data_generating_model,
        np_rng,
        FLAGS.num_features,
        FLAGS.num_samples,
        FLAGS.sigma_noise,
    )

    model = Model(tf_rng, FLAGS.num_features)
    logging.debug(model.model)

    optimizer = tf.optimizers.SGD(learning_rate=FLAGS.learning_rate)

    bar = trange(FLAGS.num_iters)
    for i in bar:
        with tf.GradientTape() as tape:
            x, y = data.get_batch(np_rng, FLAGS.batch_size)
            y_hat = model(x)
            loss = 0.5 * tf.reduce_mean(
                (y_hat - y) ** 2
            )  # implement loss function to minimize

```

Sep 07, 22 23:42

main.py

Page 4/4

```

        grads = tape.gradient(loss, model.trainable_variables)
        optimizer.apply_gradients(zip(grads, model.trainable_variables))

        bar.set_description(f"Loss @ {i} => {loss.numpy():0.6f}")
        bar.refresh()

    logging.debug(model.model)

    # print out true values versus estimates
    print("w, w_hat")
    """
compare_linear_models(data.model, model.model)
"""

    logging.info(f"Mu: {model.m}")

    fig, ax = plt.subplots(1, 2, figsize=(10, 3), dpi=200)
    ax[0].set_title("Linear Regression on Sinewave")
    ax[0].set_xlabel("x")
    h = ax[0].set_ylabel("y", labelpad=10)
    h.set_rotation(0)

    # plot with noise
    xs = np.linspace(-1.0, 2, 1000)
    xs = xs[:, np.newaxis]
    yhat = model(xs)
    ax[0].plot(xs, np.squeeze(yhat), "--", color="green")
    ax[0].plot(np.squeeze(data.x), data.y, "o", color="blue")
    ax[0].set_ylim(np.amin(data.y) * 1.5, np.amax(data.y) * 1.5)
    ax[0].set_xlim(0.1, 0.9)

    # plot sine
    true_y = np.sin(2 * np.pi * xs)
    ax[0].plot(xs, true_y, color="red")
    ax[0].legend(["Regression Estimate", "Data", "True Sine"], fontsize=9)

    # plot bases
    ax[1].set_title("Basis Functions")
    ax[1].set_xlabel("x")
    h = ax[1].set_ylabel("y", labelpad=10)
    h.set_rotation(0)

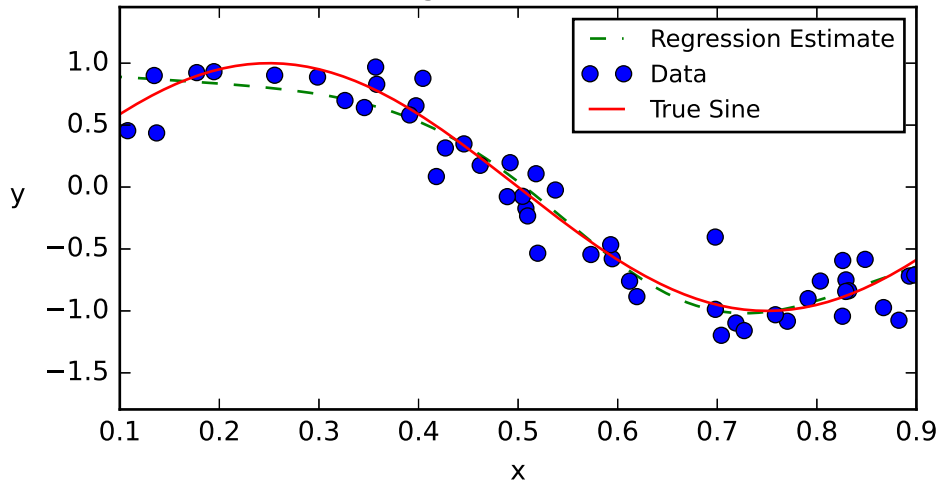
    # plot gaussians
    gaussians = np.zeros((1000, model.num_features))
    for j in range(model.num_features):
        gaussians[:, j] = tf.math.exp(
            -1 * ((xs.T - model.m[j]) ** 2) / (model.s[j] ** 2)
        )
    ax[1].plot(xs, gaussians)

    plt.tight_layout()
    plt.savefig(f"{script_path}/fit.pdf")

if __name__ == "__main__":
    app.run(main)

```

Linear Regression on Sinewave



Basis Functions

