

Sep 15, 22 0:06	hw2.py	Page 1/4
<pre>#!/bin/env python3.8 """ Simon Yoon ECE471 Deep Learning Professor Curro I had originally informed by activation function, learning rate, regularization choices by using different configurations on the TensorFlow NN playground by Carnegie Mellon students, but it did not work as intended. I had my first few runs with the layers set up as [(2,100),(100,100),(100,1)] with ELU. tf.nn.sigmoid_cross_entropy_with_logits did not work as intended. It improved when changed to RELU and when more neurons were added to the hidden layers 100-->200. Sigmoid as final activation function to keep output between [0,1]. My model was stuck at NaN for loss for nonzero lambda. Running when lambda = 0, fixed the NaN loss, but changing the lambda and learning rate to much smaller numbers helped it run with nonzero lambda. Loss was blowing up with more variables and had to divide by len(model.trainable_variables) to suppress loss. """ import os import logging import matplotlib.pyplot as plt import numpy as np from numpy import pi import tensorflow as tf from absl import app from absl import flags from tqdm import trange from dataclasses import dataclass, field, InitVar script_path = os.path.dirname(os.path.realpath(__file__)) FLAGS = flags.FLAGS flags.DEFINE_integer("num_features", 2, "Number of features in record") # spiral count flags.DEFINE_integer("num_samples", 1000, "Number of samples in dataset") flags.DEFINE_integer("batch_size", 32, "Number of samples in batch") flags.DEFINE_integer("num_iters", 1000, "Number of iterations") flags.DEFINE_float("learning_rate", 0.003, "Learning rate") flags.DEFINE_integer("random_seed", 31415, "Random seed") flags.DEFINE_float("sigma_noise", 0.1, "Standard deviation of noise random variable") flags.DEFINE_float("LAMBDA", 0.001, "Lambda") flags.DEFINE_list("layers", [(2, 200), (200, 100), (100, 1)], "Layers") flags.DEFINE_bool("debug", False, "Set logging level to debug") THETA_MAX = 4 * pi @dataclass class SpiralModel: x: np.ndarray y: np.ndarray theta: np.ndarray true: np.ndarray @dataclass class Data: model: SpiralModel rng: InitVar[np.random.Generator] num_features: int num_samples: int</pre>		

Sep 15, 22 0:06	hw2.py	Page 2/4
<pre>sigma: float x: np.ndarray = field(init=False) y: np.ndarray = field(init=False) theta: np.ndarray = field(init=False) # true: np.ndarray = field(init=False) def __post_init__(self, rng): self.index = np.arange(self.num_samples * self.num_features) self.theta = rng.uniform(0, THETA_MAX, size=(self.num_samples, 1)) self.x = np.zeros((self.num_samples * self.num_features, 2), dtype="float32") self.y = np.zeros(self.num_samples * self.num_features, dtype="float32") # spirals # formatting for clean output for c in range(self.num_features): i = range(self.num_samples * c, self.num_samples * (c + 1)) r = np.linspace(1, 15, self.num_samples) self.theta = (np.linspace(c * 3, (c + 4) * 3, self.num_samples) + np.random.randn(self.num_samples) * self.sigma) self.x[i] = np.c_[r * np.sin(self.theta), r * np.cos(self.theta)] self.y[i] = c def get_batch(self, rng, batch_size): """ Select random subset of examples for training batch """ choices = rng.choice(self.index, size=batch_size) return self.x[choices], self.y[choices] # https://tensorflow.org/guide/core/mlp_core#multilayer_perceptron_mlp_overview def xavier_init(shape): # Computes the xavier initialization values for a weight matrix in_dim, out_dim = shape xavier_lim = tf.sqrt(6.0) / tf.sqrt(tf.cast(in_dim + out_dim, tf.float32)) weight_vals = tf.random.uniform(shape=(in_dim, out_dim), minval=-xavier_lim, maxval=xavier_lim, seed=22) return weight_vals # https://www.tensorflow.org/api_docs/python/tf/Module class DenseLayer(tf.Module): def __init__(self, out_dim, weight_init=xavier_init, activation=tf.identity): : # Initialize the dimensions and activation functions self.out_dim = out_dim self.weight_init = weight_init self.activation = activation self.built = False def __call__(self, x): if not self.built: # Infer the input dimension based on first call self.in_dim = x.shape[1] # Initialize the weights and biases using Xavier scheme self.w = tf.Variable(xavier_init(shape=(self.in_dim, self.out_dim))) self.b = tf.Variable(tf.zeros(shape=(self.out_dim,)))</pre>		

Sep 15, 22 0:06

hw2.py

Page 3/4

```

        self.built = True
        # Compute the forward pass
        # z = tf.add(tf.matmul(x, self.w), self.b)
        z = x @ self.w + self.b
        return self.activation(z)

# https://www.tensorflow.org/api_docs/python/tf/Module
class MLP(tf.Module):
    def __init__(self, layers):
        self.layers = layers

    @tf.function
    def __call__(self, x, rng):
        # Execute the model's layers sequentially
        for layer in self.layers:
            x = layer(x)
        return tf.squeeze(x)

    @property
    def model(self):
        return self.layers

# https://www.tensorflow.org/guide/keras/custom_layers_and_models
def main(a):
    logging.basicConfig()

    if FLAGS.debug:
        logging.getLogger().setLevel(logging.DEBUG)

    # Safe np and tf PRNG
    seed_sequence = np.random.SeedSequence(FLAGS.random_seed)
    np_seed, tf_seed = seed_sequence.spawn(2)
    np_rng = np.random.default_rng(np_seed)
    tf_rng = tf.random.Generator.from_seed(tf_seed.entropy())

    data_generating_model = SpiralModel(
        theta=np_rng.integers(low=0, high=THETA_MAX, size=(FLAGS.num_samples, 1))
    ),
    x=np_rng.integers(low=0, high=10, size=(FLAGS.num_samples, 1)),
    y=np_rng.integers(low=0, high=10, size=(FLAGS.num_samples, 1)),
    true=np_rng.integers(low=0, high=1, size=(FLAGS.num_samples, 1)),
    )

    data = Data(
        data_generating_model,
        np_rng,
        FLAGS.num_features,
        FLAGS.num_samples,
        FLAGS.sigma_noise,
    )

    model = MLP(
        [
            DenseLayer(200, activation=tf.nn.relu),
            DenseLayer(100, activation=tf.nn.relu),
            DenseLayer(1, activation=tf.math.sigmoid),
        ]
    )

# https://ai.stackexchange.com/questions/18206/what-kind-of-optimizer-is-sug

```

Sep 15, 22 0:06

hw2.py

Page 4/4

```

gested-to-use-for-binary-classification-of-similar
optimizer = tf.keras.optimizers.Adam(learning_rate=FLAGS.learning_rate)

def cross_entropy(y, y_hat):
    loss = -y * tf.math.log(y_hat) - (1 - y) * tf.math.log(1 - y_hat)
    return tf.reduce_mean(loss)

bar = trange(FLAGS.num_iters)
for i in bar:
    with tf.GradientTape() as tape:
        x, y = data.get_batch(np_rng, FLAGS.batch_size)
        y_hat = model(x, tf_rng)
        # loss = tf.reduce_mean(
        #     tf.nn.sigmoid_cross_entropy_with_logits(logits=y_hat, labels=y
        # )
        # )
        loss = cross_entropy(y, y_hat)
        l2_reg = 0
        for var in model.trainable_variables:
            l2_reg += tf.nn.l2_loss(var)
        loss = tf.reduce_mean(
            loss + FLAGS.LAMBDA * l2_reg / len(model.trainable_variables)
        )

        grads = tape.gradient(loss, model.trainable_variables)
        optimizer.apply_gradients(zip(grads, model.trainable_variables))
        bar.set_description(f"Loss @ {i} => {loss.numpy():0.6f}")
        bar.refresh()

# logging.debug(model.model)

n = 100
x = np.linspace(-15, 15, n, dtype="float32")
xx, yy = np.meshgrid(x, x)

# reshape to fit mat dim of Model
decision = np.zeros((n, n))
for i in range(n):
    for j in range(n):
        xy = np.reshape([x[j], x[i]], (1, 2))
        p = model(xy, tf_rng)
        if p >= 0.5:
            decision[i, j] = 1

cm = plt.cm.RdBu

plt.contourf(xx, yy, decision, cmap=cm, alpha=0.5)
plt.scatter(
    data.x[:, 0], data.x[:, 1], c=data.y, cmap=cm, alpha=1, edgecolors="black"
)

plt.show()
plt.tight_layout()
plt.savefig(f"{script_path}/spiral.pdf")

if __name__ == "__main__":
    app.run(main)

```

